



Spark Training 101
Melbourne
Australia
April 2015 @ Telstra

Ned Shawa
Mark Moloney
Tim Findley
Kon

Agenda


- Spark 101 (Ned) 6:00 – 7:00
- Introduction to Scala (Mark) 7:00-7:30
- Eclipse and Spark (Tim) 7:30 – 8:00
- IntelliJ and Spark (Kon) 8:00- 8:30

Demo Prep Work

- USB will rotate across all of you for copying the USB folder
- You will run spark from your USB folder
- If you prefer Linux and you are running windows we have a vmware and a virtual box appliance that you can run
- You will have a copy of the slides in PDF in the USB folder for following commands

Basics

A Brief History: *Functional Programming for Big Data*



The need for Spark

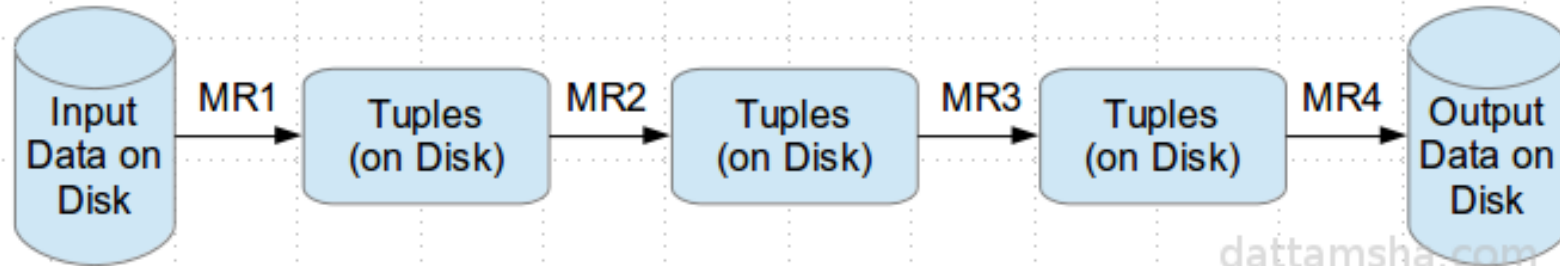
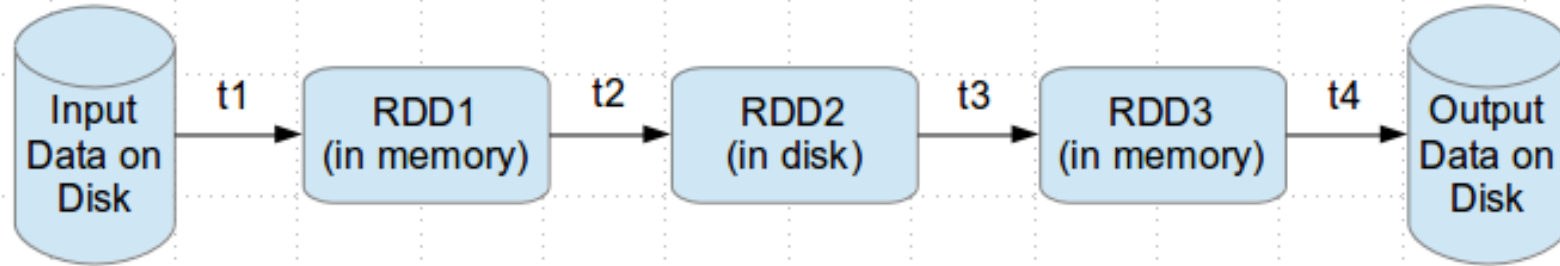
- Data movement across disks is expensive
- Apps and data reload results in re-fetching the dataset everytime
- Caching is limited
- Different data access layers are required for each usecase:
 - SQL(Hive,Hawq,Impala,..etc)
 - Graph
 - MR
 - Machine Learning

Solution

Resilient Distributed Datasets (RDDs)

Allow Apps to keep working sets in memory for efficient reuse Retain the attractive properties of MapReduce » Fault tolerance, data locality, scalability and Support a wide range of applications

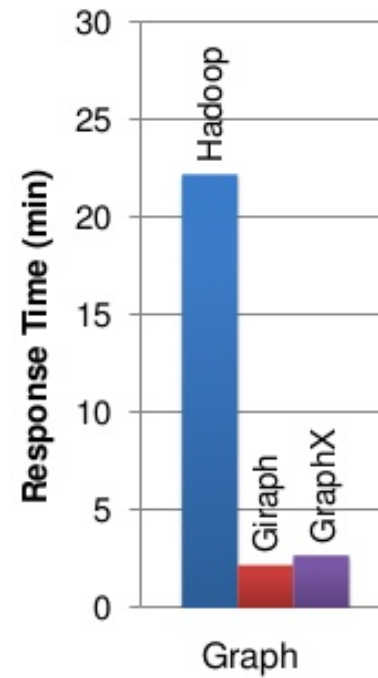


And this means?



dattamsha.com

And also means

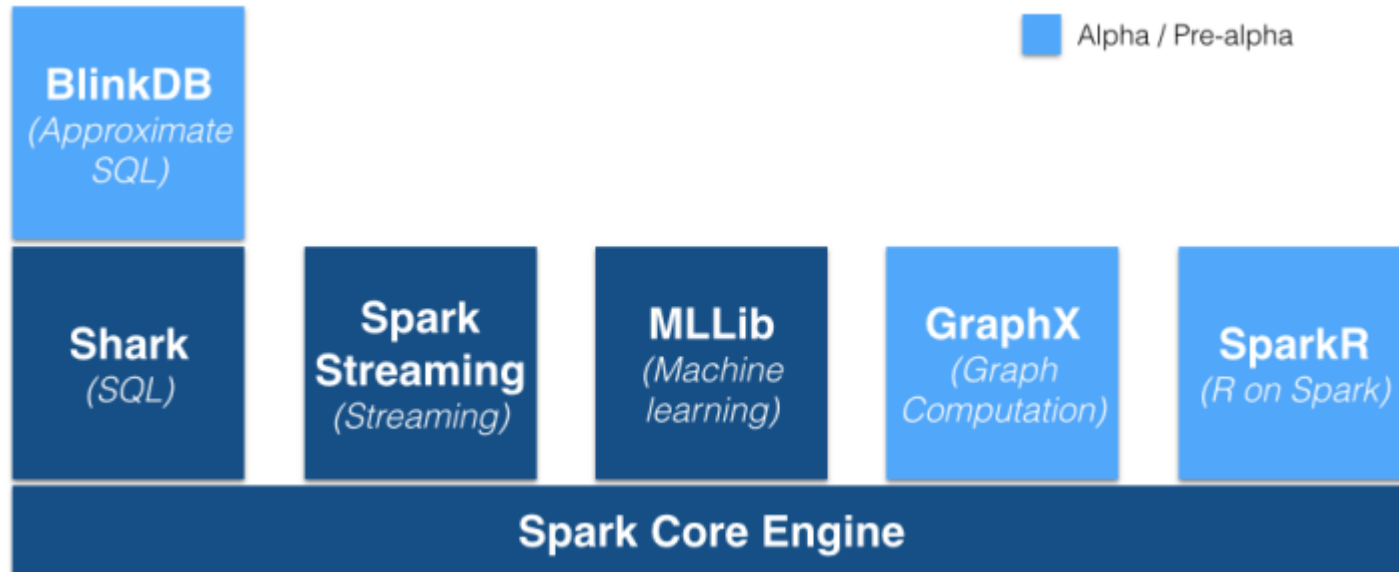
Spark Performance



Spark Programming Model

- Key idea: *resilient distributed datasets (RDDs)*
 - Distributed collections of objects that can be cached in memory across cluster nodes
 - Manipulated through various parallel operators
 - Automatically rebuilt on failure
- Interface
 - Clean language-integrated API in Scala
 - Can be used *interactively* from Scala console

Spark Components (Updated)



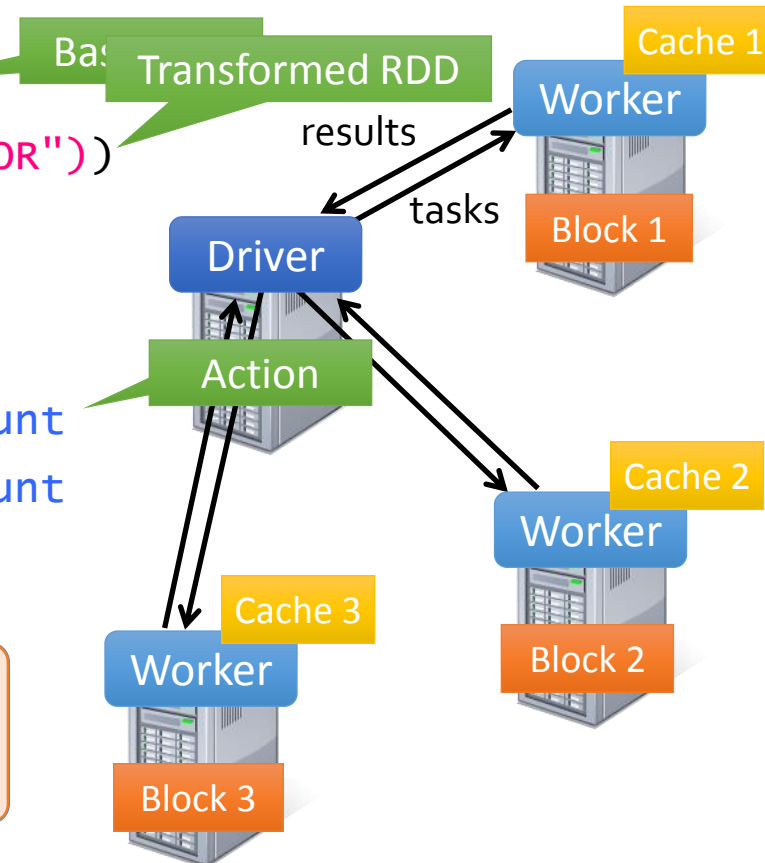
Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(_.startsWith("ERROR"))
messages = errors.map(_.split('\t')(2))
cachedMsgs = messages.cache()
```

```
cachedMsgs.filter(_.contains("foo")).count
cachedMsgs.filter(_.contains("bar")).count
. . .
```

Result: scaled to 1 TB data in 5-7 sec
(vs 170 sec for on-disk data)

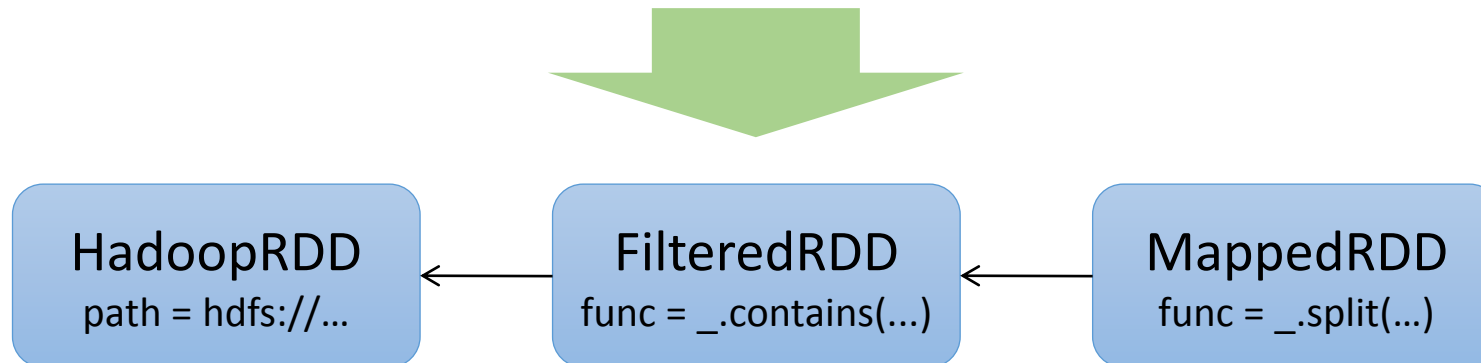


Fault Tolerance

RDDs track the series of transformations used to build them (their *lineage*) to recompute lost data

E.g:

```
messages = textFile(...).filter(_.contains("error"))  
                        .map(_.split('\t')(2))
```



Example: Logistic Regression

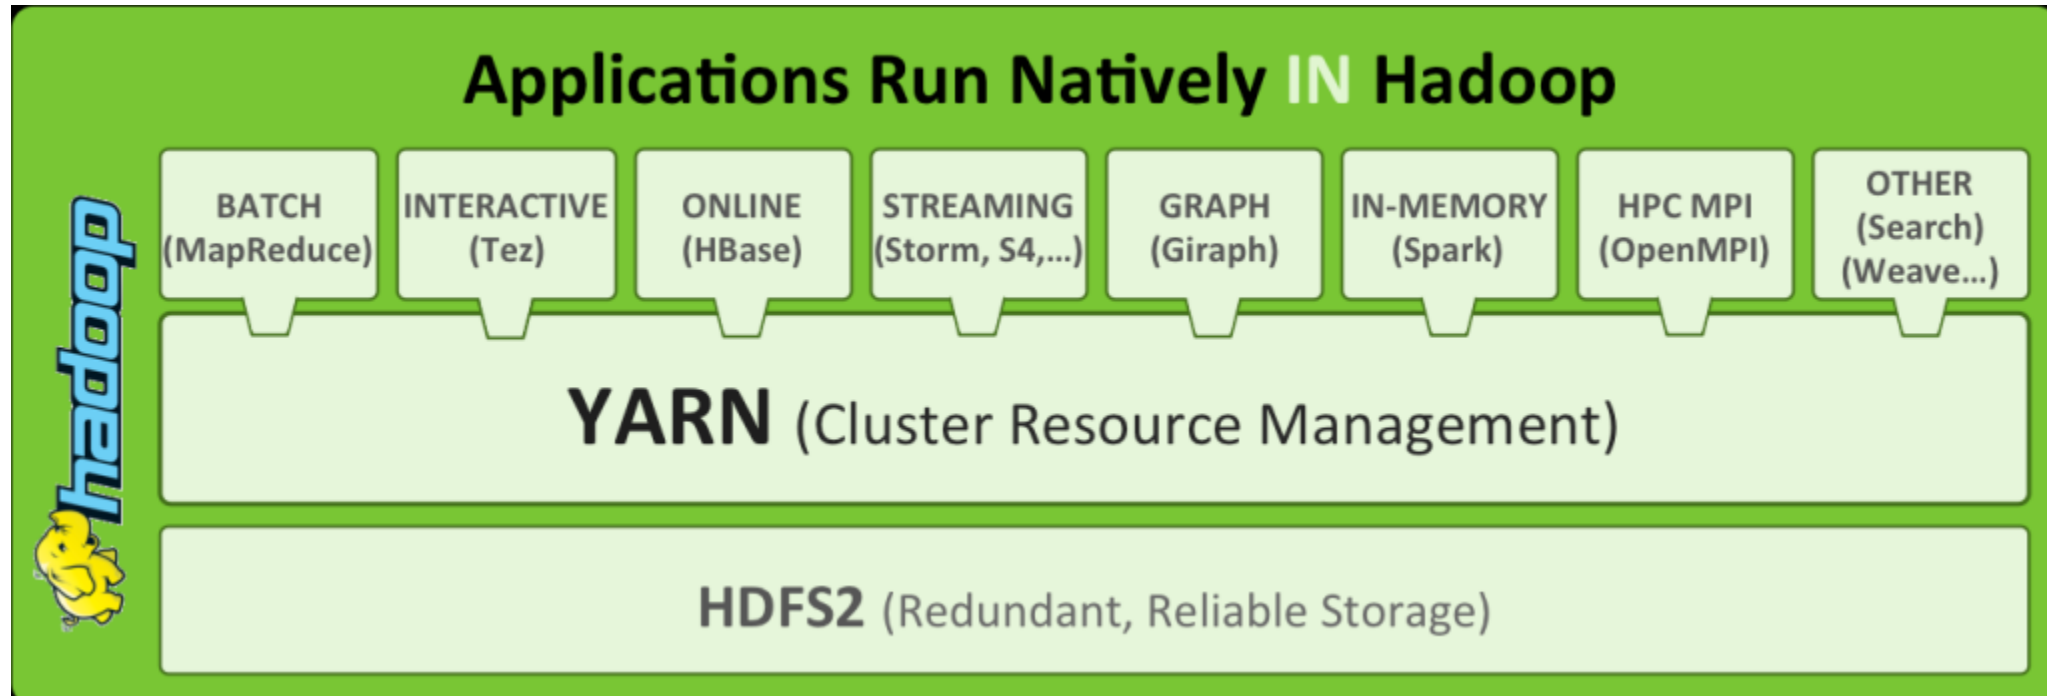
```
val data = spark.textFile(...).map(readPoint).cache()
var w = Vector.random(D)
for (i <- 1 to ITERATIONS) {
  val gradient = data.map(p =>
    (1 / (1 + exp(-p.y*(w dot p.x))) - 1) * p.y * p.x
  ).reduce(_ + _)
  w -= gradient
}
println("Final w: " + w)
```

Initial parameter vector

Load data in memory once


Repeated MapReduce steps to do gradient descent

YARN



Spark Deployment

- Standalone (Mac/Linux/Windows...etc)
- Hadoop / Yarn (Pivotal,Hortonworks..etc)
- Mesos (Cluster Computing)

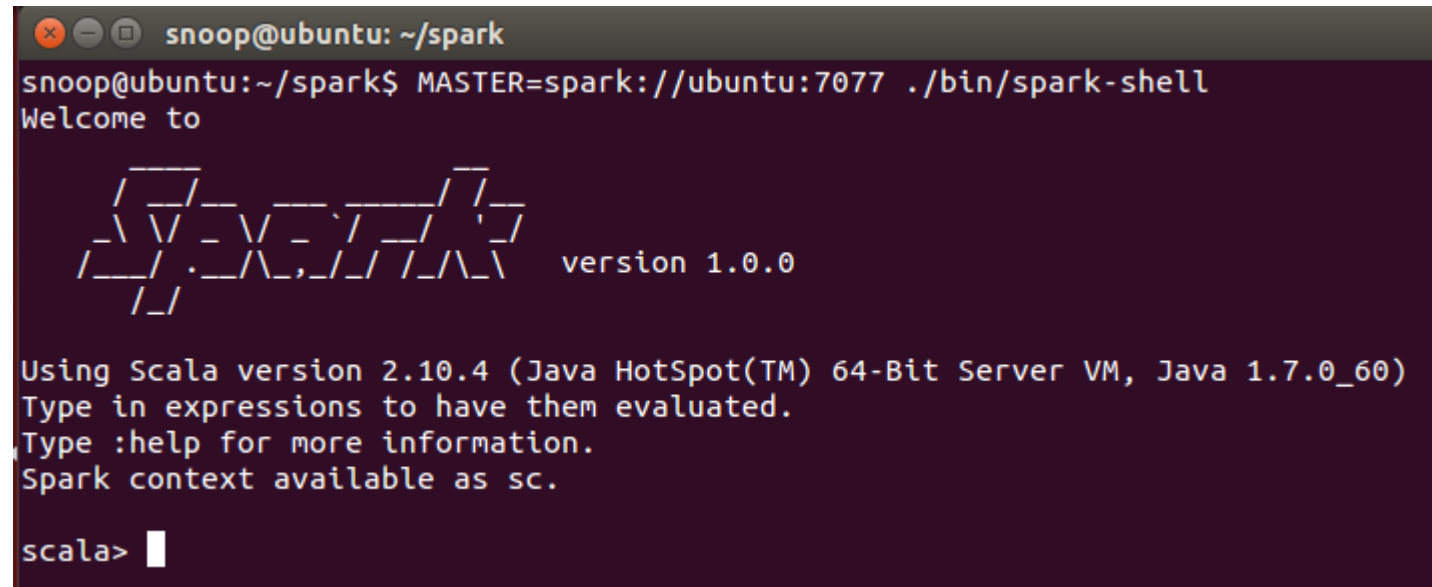


Building Spark

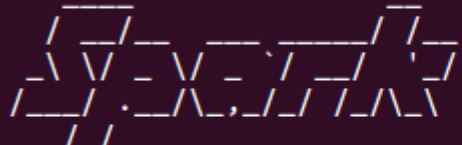
- Maven, Gradle Or SBT
- Options depends on target:
 - Hive Thrift Server
 - Hadoop Version
 - ...etc
- Example: ***mvn -Pyarn -Phadoop-2.4 -Dhadoop.version=2.4.0 -DskipTests clean package***
- Refer to <https://spark.apache.org/docs/latest/building-spark.html>

Running Spark

- Interactive Shell(/bin/spark-shell):
 - Python
 - Scala
- Built in Apps
 - Java
 - Scala
 - Python



```
snoop@ubuntu: ~/spark
snoop@ubuntu:~/spark$ MASTER=spark://ubuntu:7077 ./bin/spark-shell
Welcome to

 version 1.0.0

Using Scala version 2.10.4 (Java HotSpot(TM) 64-Bit Server VM, Java 1.7.0_60)
Type in expressions to have them evaluated.
Type :help for more information.
Spark context available as sc.

scala> 
```

Spark UI

Spark Master at spark://127.0.0.1:7077

URL: spark://127.0.0.1:7077

Workers: 1

Cores: 6 Total, 6 Used

Memory: 8.5 GB Total, 512.0 MB Used

Applications: 1 Running, 1 Completed

Drivers: 0 Running, 0 Completed

1

Workers

2

Id	Address	State	Cores	Memory
worker-20141001184107-127.0.0.1-55410	127.0.0.1:55410	ALIVE	6 (6 Used)	8.5 GB (512.0 MB Used)

Running Applications

3

ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
app-20141001184908-0000	Spark shell	6	512.0 MB	2014/10/01 18:49:08	russellspitzer	RUNNING	52 min

Completed Applications

4

ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
app-20141001184917-0001	Spark shell	0	512.0 MB	2014/10/01 18:49:17	russellspitzer	FINISHED	52 min

Spark Context (SC)

- Represents the connection to the cluster
- One SC is only allowed per JVM
- Allows you to interact and create RDDs

```
Val file = sc.textFile("some dataset")....
```

Exploring Spark

Loading dataset

```
val water_data= sc.textFile("/usb/data/vic_water.csv")
```

Basic operations

```
water_data.count()
```

```
water_data.first()
```

```
water_data.last()....
```

Exploring Spark

Filter

```
val coburg = water_data.filter(line=>line.contains("Coburg"))  
coburg.collect()
```

Caching

```
water_data.persist()  
water_data.unpersist()
```

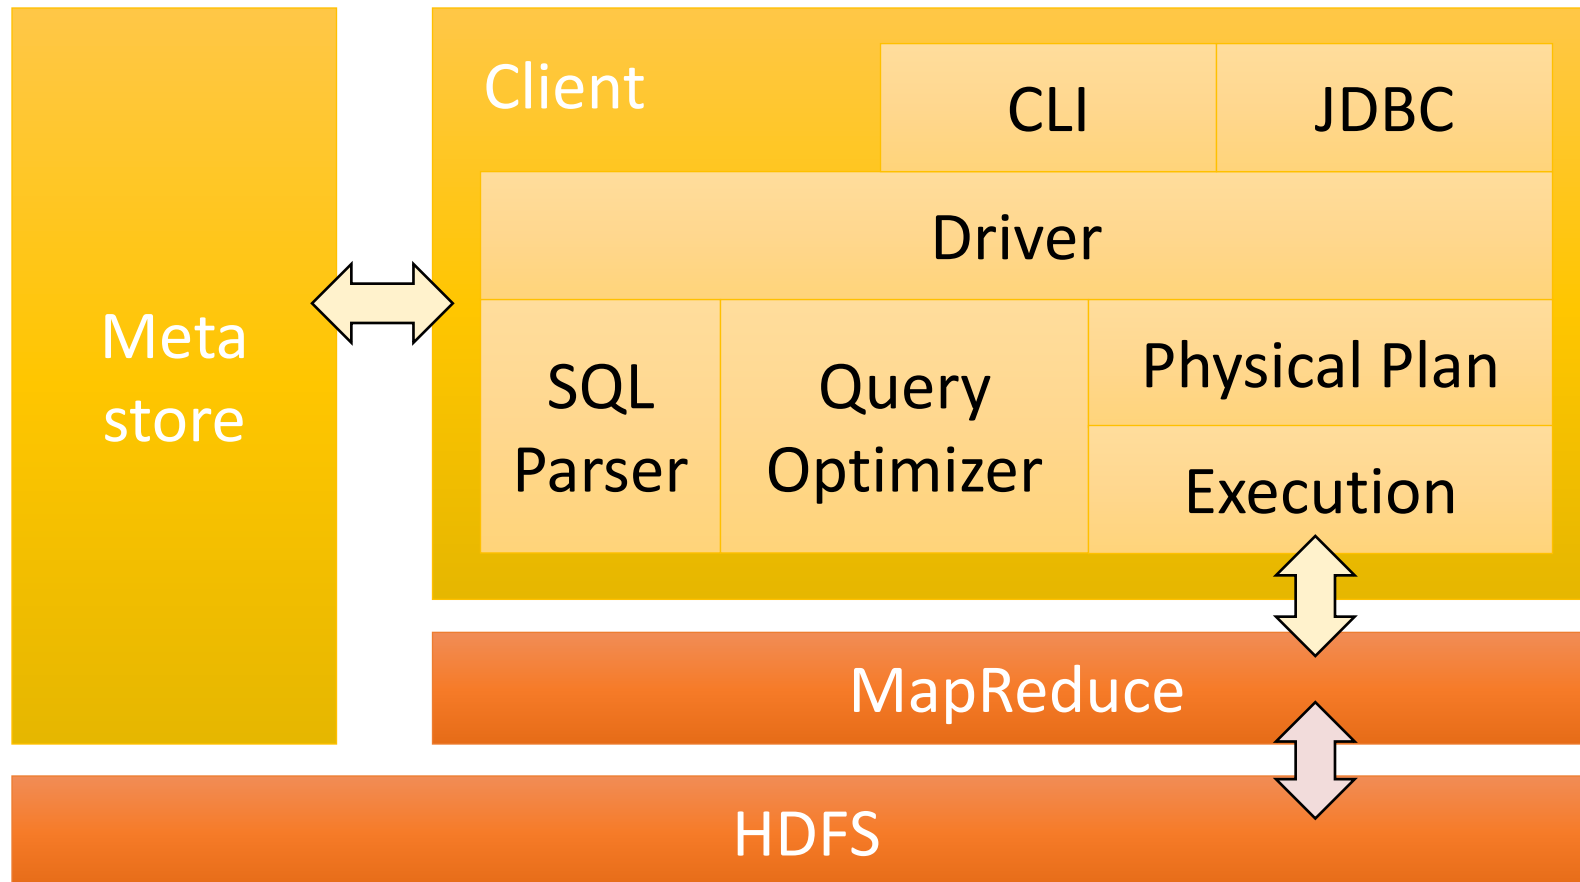
Spark SQL

- Subset of HiveQL
- Uses SchemaRDD
- Direct integration with Hive
- Many sources (JDBC,Text,JSON,Parquet...etc)

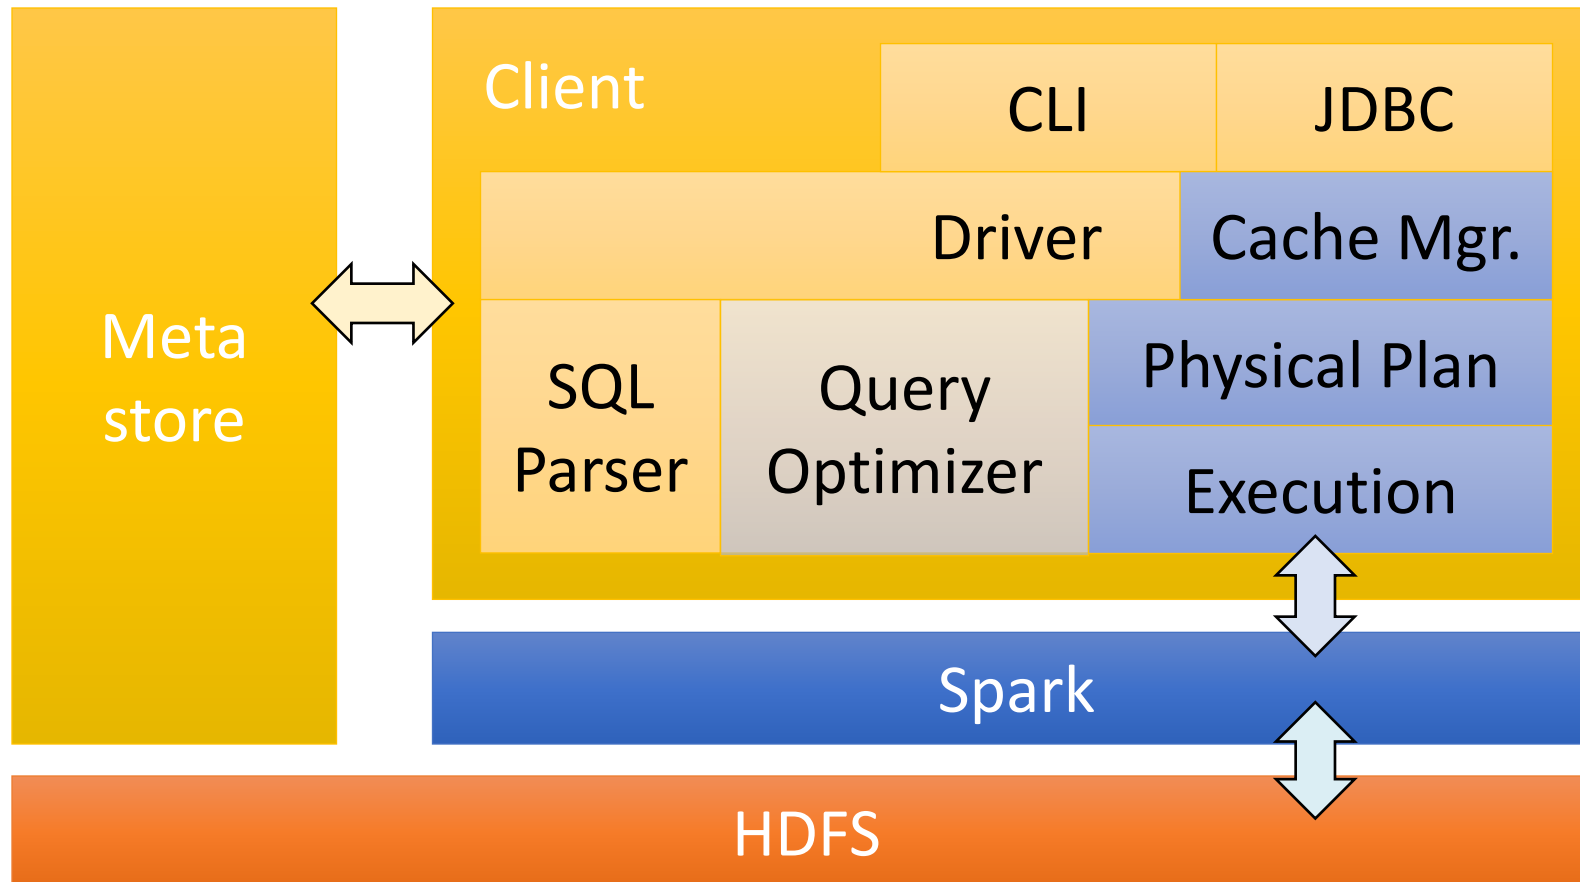
Motivation

- Hive is great, but Hadoop's execution engine makes even the smallest queries take minutes
- Scala is good for programmers, but many data users only know SQL
- **Can we extend Hive to run on Spark?**

Hive Architecture



Spark –SQL Architecture



Efficient In-Memory Storage

- Simply caching Hive records as Java objects is inefficient due to high per-object overhead
- Instead, Spark-SQL employs column-oriented storage using **arrays of primitive types**

Row Storage

1	john	4.1
2	mike	3.5
3	sally	6.4

Column Storage

1	2	3
john	mike	sally
4.1	3.5	6.4

Efficient In-Memory Storage

- Simply caching Hive records as Java objects is inefficient due to high per-object overhead
- Instead, Spark-SQL employs column-oriented storage using **arrays of primitive types**

Row Storage

Column Storage

Benefit: similarly compact size to serialized data,
but >5x faster to access

3	sally	6.4
---	-------	-----

4.1	3.5	6.4
-----	-----	-----

Spark SQL (1.3)

Built-In

{ JSON }



JDBC

Parquet



MySQL



amazon web services S3



External



dBase

APACHE
HBASE

elasticsearch.



amazon
web services
Amazon Redshift

and more...

What are DataFrames?

- Distributed Collection of Data organized in Columns
- Equivalent to Tables in Databases or DataFrame in R/PYTHON
- Much richer optimization than any other implementation of DF
- Can be constructed from a wide variety of sources and APIs

Exploring Spark SQL

```
import sqlContext.implicits._
```

```
case class water (postcode: Int, suburb:String, cons_08: Int, cons_09: Int)
```

```
val water_table
```

```
=water_data.map(_.split(" | ")).map(t=>water(t(0).trim.toInt,t(1),t(2).trim.toInt,t(3).trim.toInt)).toDF()
```

```
water_table.registerTempTable("water")
```

Reading/Writing a DataFrame

```
val df = sqlContext.load("/usb/data/people.json", "json")
Val df =sqlContext.load("/usb/data/users.parquet", "parquet")
df.show()
df.printSchema()
df.select ("name").show()
df.select("name","favorite_color").show()
df.select("name").save("/usb/data/names.parquet", "parquet")
df.registerTempTable("df")
sqlContext.sql("select * from df").foreach(println)
```


A Brief Introduction to Scala (and how it can help you build better data apps)

Mark Moloney
markmo@me.com

The moot

- Functional programming style works well with data-driven applications
- Scala supports a functional programming style without losing integration options and operational support of the Java ecosystem
- Knowing Scala == Spark power user

Some facts and history

- Scala is an **object-functional** language
- Object-oriented, C-style syntax (but with less boilerplate)
- Has features of FP languages (e.g. Haskell, Scheme)
- Statically typed but uses type inference (can sometimes appear dynamically typed)
- A key focus of the language has been on making development of concurrent systems easier.
- Scala code compiles to Java bytecode to leverage the Java Virtual Machine (JVM)
 - 20 years of tuning to perform comparably to C, plus a wealth of libraries, management tools, and acceptance by most ops groups of large companies
 - You can create Java classes and call Java methods directly from Scala
- It was created by Martin Odersky (who had previously worked on the Java compiler) and publically released in early 2004
- Commercial support for the language is provided by Typesafe, a company founded by Odersky

What is Functional Programming anyway?

- Imperative programming has been the norm:
 - modifying mutable variables (meaning data that can change after initialization)
 - using assignments
 - and control structures such as if-then-else, loops, etc.
- However, “Concurrency is the Dr. Evil to mutable state’s Austin Powers.”¹
- If two different threads can change the same data at the same time, its difficult to guarantee that the execution will leave the data in a valid state
- In contrast, functional programming favours:
 - immutable values
 - functions that always return a value, and if given the same inputs, will always return the same value
 - recursion and “flow syntax”

Basics

```
val a = "Can't touch this"  
a = "We'll see"
```

```
var b = "Change me"  
b = "OK"
```

```
def square(x:Int) = x * x
```

```
def sumOfSquares(x:Int, y:Int): Int = {  
    val x2 = x * x  
    val y2 = y * y  
    x2 + y2  
}
```

```
val numbers = List(1, 2, 3, 4)  
numbers(2)  
val uniqueNumbers = Set(1, 1, 2)  
val hostPort = ("localhost", 80)  
hostPort._1  
hostPort._2  
val fooMap = Map("foo" -> "bar", "fi" -> "baz")  
val v = fooMap.get("foo")  
v  
v.get  
fooMap.get("binkle")  
fooMap.getOrElse("binkle", "boo")
```

```
val range = 0 until 10
```

Type Inference

Scala

```
val b = 1 < 2 // x is a Boolean
```

```
val word = "Hello" // word is a String
```

Java

```
Boolean b = 1 < 2
```

```
String word = "Hello"
```

Functional Style

Scala

```
val numbers = List(1, 2, 3, 4)

numbers.map((x: Int) => x * 2)
numbers.map(_ * 2)

val result = numbers.map(_ * 2)
result.foreach((x: Int) => {
    println(x)
})
result.foreach(println)

numbers.map(_ * 2).filter(_ < 5)

result.take(1)

numbers.reduce((x: Int, y: Int) => x + y)
numbers.reduce(_ + _)
def plus(x: Int, y: Int) = x + y
numbers.reduce(plus)

numbers.foldLeft(0) { (m: Int, n: Int) => {
    println("m: " + m + " n: " + n)
    m + n
}

val nestedNumbers = List(List(1, 2), List(3, 4))
nestedNumbers.map(xs => xs.map(_ * 2))
nestedNumbers.flatMap(xs => xs.map(_ * 2))
```

Java

```
int[] numbers = new int[] { 1, 2, 3, 4 };
//or
List<Integer> numbers = new ArrayList<Integer>() {{
    add(1);
    add(2);
    add(3);
    add(4);
}};

// map equivalent
int[] mapResult = new int[numbers.length];
for (int i = 0; i < numbers.length; i++) {
    mapResult[i] = numbers[i] * 2;
}

// reduce equivalent
int total = 0;
for (int i : numbers) {
    total += i;
}
```

Case Classes

Scala

```
case class Person(firstName: String, lastName: String)
```

```
val fred = Person("fred", "wilson")
```

```
println(fred.lastName)  
> Wilson
```

And supports pattern matching!

```
def greetFred(randomPerson: Any) = {  
  randomPerson match {  
    case Person("fred", _) => println("G'day Fred")  
    case _                  => println("Hello")  
  }  
}
```

```
greetFred(fred)  
> G'day Fred
```

Java



```
public class Person implements java.io.Serializable {  
  
    private String firstName;  
    private String lastName;  
  
    public String getFirstName() {  
        return firstName;  
    }  
  
    public void setFirstName(String firstName) {  
        this.firstName = firstName;  
    }  
  
    public String getLastName() {  
        return lastName;  
    }  
  
    public void setLastName(String lastName) {  
        this.lastName = lastName;  
    }  
  
    @Override  
    public boolean equals(Object o) {  
        if (this == o) return true;  
        if (o == null || getClass() != o.getClass()) return false;  
  
        Person person = (Person) o;  
  
        if (firstName != null ? !firstName.equals(person.firstName) : person.firstName != null) return false;  
        if (lastName != null ? !lastName.equals(person.lastName) : person.lastName != null) return false;  
  
        return true;  
    }  
  
    @Override  
    public int hashCode() {  
        int result = firstName != null ? firstName.hashCode() : 0;  
        result = 31 * result + (lastName != null ? lastName.hashCode() : 0);  
        return result;  
    }  
  
    @Override  
    protected Object clone() throws CloneNotSupportedException {  
        return super.clone();  
    }  
  
    @Override  
    public String toString() {  
        return "Person{" +  
            "firstName='" + firstName + '\'' +  
            ", lastName='" + lastName + '\'' +  
            "}";  
    }  
}
```


Pattern Matching

This example reads n-gram counts from a file:

```
9 WORDTAG 0 Test
11 WORDTAG 0 cysts
43 WORDTAG 0 splice
6 WORDTAG 0 extensively
1 WORDTAG I-GENE heterodimer
```

```
13796 2-GRAM * *
749 3-GRAM * * I-GENE
11320 3-GRAM I-GENE 0 0
9622 3-GRAM I-GENE I-GENE 0
```

```
def extractCounts() {
  var counts: Source = null
  try {
    counts = Source.fromFile("/data/h1-p/gene.counts")
    counts.getLines().foreach(_ split "\\s+" match {
      case Array(k, "WORDTAG", tag, word)      => { wordTagCounts((tag, word)) = k.toInt; }
      case Array(k, "1-GRAM", tag)              => { unigramCounts(tag) = k.toInt }
      case Array(k, "2-GRAM", tag1, tag2)       => { bigramCounts((tag1, tag2)) = k.toInt }
      case Array(k, "3-GRAM", tag1, tag2, tag3) => { trigramCounts((tag1, tag2, tag3)) = k.toInt }
    })
  } finally {
    if (counts != null) counts.close()
  }
}
```

Extras if time

Working with non-tabular data

Scala

XML

```
val movies =  
  <movies>  
    <movie genre="action">Pirates of the Carribean</movie>  
    <movie genre="fairytale">Edward Scissorhands</movie>  
  </movies>  
  
movies \ "movie"  
(movies \ "movie")(0) \ "@genre"
```

JSON

```
import play.api.libs.json._  
  
case class Person(name: String, age: Int)  
  
val json: JsValue = Json.parse("""  
{  
  "name": "Billy Bob",  
  "age": 42  
}  
""")  
  
val name = json \ "name"  
  
implicit val personReads: Reads[Person] = (  
  (JsPath \ "name").read[String] and  
  (JsPath \ "age").read[Int]  
) (Person.apply _)  
  
val personResult: JsResult[Person] = json.validate[Person]  
val person = personResult.get
```

Isn't creating new values inefficient?

- The immutable data structures in Scala are designed to be efficient. For example:

Consider an immutable singly-linked list:

$L = d \rightarrow c \rightarrow b \rightarrow a$



L can be stored with just a reference to the head of the list, the element d

Suppose we want to create a new list L2 with element **e** added to the head of the list

$L2 = e \rightarrow (\text{pointer to } d)$

However, suppose we want to create a new list L3 with the elements $d \rightarrow c \rightarrow \mathbf{j} \rightarrow b \rightarrow a$

We can't mutate L since that would also change L2, therefore we need to copy elements d, c

$L3 = d \rightarrow c \rightarrow j \rightarrow (\text{pointer to } b)$

Similar to how Git manages branches, commits up to a common ancestor are not copied

What is a Map Operation?

Using JavaScript as an example:

```
function square(a) {  
    for (i = 0; i < a.length; i++) {  
        a[i] = a[i]^2;  
    }  
}
```

Generalized Map Operation

```
function map(fn, a) {  
    for (i = 0; i < a.length; i++) {  
        a[i] = fn(a[i]);  
    }  
}
```

Invoked as:

```
map(function (x) { return x^2; }, a);
```

What is a Reduce Operation?

```
function sum(a) {  
    var s = 0;  
    for (i = 0; i < a.length; i++) {  
        s += a[i];  
    }  
    return s;  
}
```

Generalized Reduce Operation

```
function reduce(fn, a, init) {  
    var s = init;  
    for (i = 0; i < a.length; i++) {  
        s = fn(s, a[i]);  
    }  
    return s;  
}
```

Invoked as:

```
reduce(function (a, b) { return a + b; }, a, 0);
```


Machine Learning Application of Map-Reduce

- Batch gradient descent:

$$\theta_j := \theta_j - \alpha \frac{1}{400} \sum_{i=1}^{400} (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

Map

- Machine 1:

$$(x^{(1)}, y^{(1)}), \dots, (x^{(100)}, y^{(100)})$$
$$temp_j^{(1)} = \sum_{i=1}^{100} (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)}$$

- Machine 2:

$$(x^{(101)}, y^{(101)}), \dots, (x^{(200)}, y^{(200)})$$
$$temp_j^{(2)} = \sum_{i=101}^{200} (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)}$$

- Machine 3:

$$(x^{(201)}, y^{(201)}), \dots, (x^{(300)}, y^{(300)})$$
$$temp_j^{(3)} = \sum_{i=201}^{300} (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)}$$

- Machine 4:

$$(x^{(301)}, y^{(301)}), \dots, (x^{(400)}, y^{(400)})$$
$$temp_j^{(4)} = \sum_{i=301}^{400} (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)}$$

Reduce

$$\theta_j := \theta_j - \alpha \frac{1}{400} (temp_j^{(1)} + temp_j^{(2)} + temp_j^{(3)} + temp_j^{(4)})$$

